

UNIVERSITY OF BRISTOL

INTRODUCTION TO MACHINE LEARNING

PROJECT REPORT

---

# Naive Bayes Spam Filter

---

*Authors:*

Prisca Aeby, Alexis Semple

December 9, 2014

## **Abstract**

With today's world becoming more and more flooded with e-mails, spammers are responding with an alarming increase of unsolicited messages. The natural reaction to this has been the development of many different spam filtering techniques, of varying efficiency and reliability. We are going to be looking more specifically at one of these techniques, Naive Bayes spam filtering, how it can be implemented and improved, and how it compares to other filtering techniques.

# Contents

1	Introduction . . . . .	1
2	The Algorithm . . . . .	2
2.1	Training . . . . .	2
2.2	Classification . . . . .	4
3	Analysis of results and improving the classifier . . . . .	5
3.1	10-fold cross-validation . . . . .	5

## 1 Introduction

Naive Bayes text classification is a popular technique used in spam e-mail filtering. It relies on correlating the use of tokens (in our case words) and then using Bayesian inference to calculate the probability that a given e-mail is spam or ham.<sup>1</sup> Naive Bayes classifiers suppose strong independence between the features of the model.

In the case of our project, the goal was to develop a binomial Naive Bayes text classifier written in Java to be tested and trained on real world data, or more specifically a large set of real world e-mails.

---

<sup>1</sup>Wikipedia on Naive Bayes spam filtering - [http://www.wikiwand.com/en/Naive\\_Bayes\\_spam\\_filtering](http://www.wikiwand.com/en/Naive_Bayes_spam_filtering)

## 2 The Algorithm

We based the main algorithm of our spam filter on `LEARN_NAIVE_BAYES_TEXT` and `CLASSIFY_NAIVE_BAYES_TEXT` found in Tom Mitchell's book (on page 183).<sup>2</sup> Our program is split into two parts.

### 2.1 Training

The file `filter.train.java` runs all of the computations and produces the results necessary to the supervised training part of the Naive Bayes text classification algorithm. In it, we iterate over all the files contained in a 'training set' directory and create a `java.util.HashMap` of all of the words therein. Each word is then a key in the `HashMap` and is associated with its value, a static array, each containing 4 precise values in the following order (`double[0..3]`):

- `double[0]`: The number of times the word has occurred in all files marked 'spam\*' in the training data.
- `double[1]`: The number of times the word has occurred in all files marked 'ham\*' in the training data.
- `double[2]`: The probability that a randomly drawn word from a document in class *Spam* will be this word  $w_k \rightarrow P(w_k|Spam)$
- `double[3]`: The probability that a randomly drawn word from a document in class *Ham* will be this word  $w_k \rightarrow P(w_k|Ham)$

where we have  $pSpam$  and  $pHam$  respectively being computed by dividing the number of files marked spam\* (or ham\*) over the total number of training files. The values of `double[2]` and `double[3]` as explained above are computed as follows:

---

```
for(String word: vocabulary.keySet()){  
    double[] tab = vocabulary.get(word);  
    tab[2] = (tab[0] + 1) / ((double)(totalWords[0] +  
        vocabulary.size()));
```

---

<sup>2</sup><http://personal.disco.unimib.it/Vanneschi/McGrawHill.-MachineLearning-Tom.Mitchell.pdf>

```
tab[3] = (tab[1] + 1) / ((double)(totalWords[1] +  
    vocabulary.size()));  
vocabulary.put(word, tab);  
}
```

---

where `totalWords[0..1]` contains the total number of words in all spam (resp. all ham) files and `vocabulary` is the `HashMap` of all words. According to the project description, we made the word selection process case-sensitive, and also decided to interpret any form of punctuation or whitespace as a word separator. So the scanner iterating over all words in file `f` behaves like such:

---

```
scanner = new Scanner(f).useDelimiter("[\\s\\p{Punct}]+");
```

---

The choice concerning punctuation was a personal one and based on the assumption that we had enough training data provided to ignore varying forms of punctuation. Results may have varied slightly had we taken these different forms into account, particularly considering use that is excessive and non-conform to standards which is dominant in some types of spam.

At the end of the program, the content of `vocabulary` is written to a text file named `trainging_data.txt` in a specific format: the first line contains *pSpam* and *pHam* in that order, separated by a space. Each of the following lines is dedicated to a single entry of the `vocabulary` `HashMap`, starting by the word itself, followed by the four values of its corresponding array, each separated by a space. The result looks something like this:

---

```
...  
invested 7.0 0.0 2.088031403992316E-5 4.980206170575049E-7  
...
```

---

meaning that for the word 'invested' we have that it appears 7 times in files marked 'spam\*', 0 times in files marked 'ham\*', has a probability of  $2.088031403992316 * 10^{-5}$  of being drawn at random from class *Spam* and a probability of  $4.980206170575049 * 10^{-7}$  of being drawn at random from class *Ham*. One obvious observation is that these probabilities are extremely small and can pose problems of arithmetic underflow later on, in the classification part of the algorithm. The solution we implemented to avoid this is to use

logarithms, which converts very small positive numbers into large negatives. So the code given above for computing `double[2]` and `double[3]` becomes

---

```
...
tab[2] = Math.log(tab[0] + 1) / (((double)(totalWords[0] +
    vocabulary.size())));
tab[3] = Math.log((tab[1] + 1) / (((double)(totalWords[1] +
    vocabulary.size())));
...
```

---

and the much more reasonable result in the text file is

---

```
...
invested 7.0 0.0 -10.776703754836734 -14.512624361064837
...
```

---

## 2.2 Classification

Once the training is completed, the program in `filter.java` can then use the results in `training_data.txt` to rebuild a Hashmap of `<word, double[]>` and classify instances provided to it. The formula provided in Mitchell's book returns an estimated target value for a given instance that is being tested. For spam

$$p_1 = p_{Spam} \prod_{w_k \in vocabulary} P(w_k | Spam)$$

and for ham

$$p_2 = p_{Ham} \prod_{w_k \in vocabulary} P(w_k | Ham)$$

and the return value is

---

```
if(max(p1, p2) == p1) System.out.println("spam\n");
else System.out.println("ham\n");
```

---

However, since in the training part we modified  $P(w_k | Spam)$  to  $\log(P(w_k | Spam))$  (resp.  $P(w_k | Ham)$  to  $\log(P(w_k | Ham))$ ) we get

$$p_1 = \log(pSpam) + \sum_{w_k \in \text{vocabulary}} \log(P(w_k|Spam))$$

with the same going for  $p_2$ , obviously. This determines whether a test instance is classified as spam (for  $p_1 > p_2$ ) or ham (otherwise).

## 3 Analysis of results and improving the classifier

### 3.1 10-fold cross-validation

As indicated by the project specification, in order to test the consistency of the classifier and to evaluate its precision, we performed a 10-fold cross-validation on the provided training set, which means that we split the training set into 10 test sets of equal size, and we did so randomly in order to obtain an approximately equal spread of members of both classes in all test sets. To each test set (10% of the original training set, i.e. 250 files) we then match the remaining 90% as a training set. So we end up being able to test our classifier on 10 different test sets with 250 test instances and different training sets paired with them.

We have provided the bash scripts that we used to perform the random splitting into 10 test-training pairs (called `divide_train.bash`) and the cross-validation (called `cross_validation.bash`). The latter outputs a result file for each test-training pair with the number of successful classifications and the total number of test instances. It also saves every misclassified instance to a separate folder for further examination.